

Geospatial models Serving and Benchmarking with vLLM

Goal: Prepare and **serve** a
TerraTorch models with **vLLM**
and **benchmark** performance

Agenda

- Step 1: Serve a TerraTorch model with vLLM
- Step 2: Benchmark the inference server with *vLLM bench* and *GuideLLM*
- Step 3: Primer on how to scale vLLM in Kubernetes
- Bonus step: Customize vLLM with IOProcessors


vLLM

vLLM is an [open-source](#) library to [high-performance inference and serving](#) of foundation models.


- Standard [REST API](#) for easy integration
- High GPU utilisation via continuous batching & scheduling
- Community-driven: [extensible, transparent](#), less lock-in
- Cloud-native ready: fits [Kubernetes](#) deployments
- Benchmarking tools: vllm bench, GuideLLM for [latency/throughput](#)




vLLM inference for EO models




vLLM inference for EO models



Storing models



Filesystem



Containers



HuggingFace

Storing models

 **Hugging Face**

Search models, datasets, users...

Models

 **mgazz** / **prithvi-eo-flood**

like 0

License: apache-2.0

Model card

Files and versions

xet

Community

Settings

main

prithvi-eo-flood

391 MB

 **mgazz**

Upload 2 files

4454e57


VERIFIED

 .gitattributes	 Safe	1.52 kB	
 README.md	 Safe	28 Bytes	
 config_deploy.yaml		4.42 kB	
 state_dict.ckpt		391 MB	 xet 

Prepare Model




model.ckpt



state_dict.bin

config_deploy.yaml



config.json

config.json

```
{  
  "architectures": [  
    "Terratorch"  
  ],  
  "pretrained_cfg": {  
    .....
```

Reference to the
Model Executor

```
    "data": {  
      "class_path": "terratorch.datamodules.GenericNonGeoSegmentationDataModule",  
      .....
```

Pre-process
Input

```
    "model": {  
      "class_path": "terratorch.tasks.SemanticSegmentationTask",  
      .....
```

Type of Model

```
    "input": {  
      "target": "pixel_values",  
      "data": {  
        "pixel_values": { "type": "torch.Tensor", "shape": [6,512,512]}  
      }  
    }  
  }  
}
```

Input validation
and vLLM
warm-up

What to expect

```
sagemaker-user@default:~/ESA-NASA-Workshop-2026/Day 2/Track 3/Operational-GeoAI-Model-Serving/Part2/1_run_model_in_vllm$ \
> ls -lah
total 746M
drwxr-xr-x. 4 sagemaker-user users 4.0K May 19 13:25 .
drwxr-xr-x. 9 sagemaker-user users 4.0K May 19 13:25 ..
drwxr-xr-x. 3 sagemaker-user users 25 May 19 13:21 .cache
drwxr-xr-x. 2 sagemaker-user users 6 May 19 13:20 .ipynb_checkpoints
-rw-r--r--. 1 sagemaker-user users 6.4K May 19 13:12 README.md
-rw-r--r--. 1 sagemaker-user users 5.9K May 19 13:25 config.json
-rw-r--r--. 1 sagemaker-user users 4.4K May 19 13:21 config_deploy.yaml
-rwxr-xr-x. 1 sagemaker-user users 3.1K May 19 13:12 convert_ckpt_to_bin.py
-rw-r--r--. 1 sagemaker-user users 559K May 19 13:12 inference.ipynb
-rw-r--r--. 1 sagemaker-user users 13K May 19 13:12 mask.tiff
-rw-r--r--. 1 sagemaker-user users 373M May 19 13:25 state_dict.bin
-rw-r--r--. 1 sagemaker-user users 373M May 19 13:21 state_dict.ckpt
-rw-r--r--. 1 sagemaker-user users 3.0K May 19 13:12 vllm_config_generator.py
```


Model Preparation Time!

(follow `1_run_model_in_vllm/README.md`)

Serve the model with vLLM

```
vllm serve 1_run_model_in_vllm \
  --skip-tokenizer-init \
  --enable-mm-embeds \
  --max-num-seqs 32 \
  --io-processor-plugin \
  terratorch_segmentation \
  --enforce-eager
```

Serve the model with vLLM

```
vllm serve 1_run_model_in_vllm \ #Model Name (path)
  --skip-tokenizer-init \ #No Tokenizer for EO
  --enable-mm-embeds \ #Turn on multi-modal
  --max-num-seqs 32 \ #Control parallelism
  --io-processor-plugin \
  terratorch_segmentation \ #IOProc Selection
  --enforce-eager #Disable CUDA graph
```


What to expect

```
299]
299]
299]  version 0.19.1
299] model 1_run_model_in_vllm
299]
299]
233] non-default args: {'model_tag': '1_run_model_in_vllm', 'model': '1_run_model_in_vllm',
e, 'io_processor_plugin': 'terratorch_segmentation', 'enable_mm_embeds': True, 'max_num_seqs
```


What to expect

299]

299]

299]

299]

299]

299]

233] non-default args: {'model_t
, 'io_processor_plugin': 'terrat



```
19 13:32:52 [api_server.py:596] Starting vLLM server on http://0.0.0.0:8000
19 13:32:52 [launcher.py:37] Available routes are:
19 13:32:52 [launcher.py:46] Route: /openapi.json, Methods: GET, HEAD
19 13:32:52 [launcher.py:46] Route: /docs, Methods: GET, HEAD
19 13:32:52 [launcher.py:46] Route: /docs/oauth2-redirect, Methods: GET, HEA
19 13:32:52 [launcher.py:46] Route: /redoc, Methods: GET, HEAD
19 13:32:52 [launcher.py:46] Route: /tokenize, Methods: POST
19 13:32:52 [launcher.py:46] Route: /detokenize, Methods: POST
19 13:32:52 [launcher.py:46] Route: /load, Methods: GET
19 13:32:52 [launcher.py:46] Route: /version, Methods: GET
19 13:32:52 [launcher.py:46] Route: /health, Methods: GET
19 13:32:52 [launcher.py:46] Route: /metrics, Methods: GET
19 13:32:52 [launcher.py:46] Route: /v1/models, Methods: GET
19 13:32:52 [launcher.py:46] Route: /ping, Methods: GET
19 13:32:52 [launcher.py:46] Route: /ping, Methods: POST
19 13:32:52 [launcher.py:46] Route: /invocations, Methods: POST
19 13:32:52 [launcher.py:46] Route: /pooling, Methods: POST
Started server process [1357]
Waiting for application startup.
Application startup complete.
```


Serve the model with vLLM

```
vllm serve mgazz/prithvi-eo-flood-vLLM \
  --skip-tokenizer-init \
  --enable-mm-embeds \
  --max-num-seqs 32 \
  --io-processor-plugin terratorch_segmentation \
  --enforce-eager
```

Model Serving and Inference Time!

(continue with the instructions in the README.md
and run the example in inference.ipynb)

Benchmark Tools

Key questions:

- **Baseline Performance**
 - What is the latency/throughput running an inference with **vLLM**?
- **Scalability behaviour**
 - How does the throughput/latency of the vLLM service change as we **increase clients requests**?
- **Production Reliability (SLO)**
 - How should the inference stack be designed to **meet latency and throughput SLOs**?
 - What configurations ensure **predictable performance under load**?

Benchmark Tools

	vLLM Bench	Guidellm
Organization	vLLM	vLLM
Location	vLLM project	Standalone project
Best for	Single-point measure Regression tests	Performance sweep Custom arrival patterns

Jsonl Dataset entries

```
{  
  "prompt": {  
    "data": {  
      "data": ".../samples/India_900498_S2Hand.tif",  
      "data_format": "path",  
      "out_data_format": "b64_json",  
      "indices": [1,2,3,8,11,12]  
    }, ...  
  }  
}
```

Jsonl Dataset entries

```
{  
  "prompt": {  
    "data": {  
      "data": ".../samples/India_900498_S2Hand.tif",  
      "data_format": "path",  
      "out_data_format": "b64_json",  
      "indices": [1,2,3,8,11,12]  
    }, ...  
  }  
}
```


vllm bench

```
vllm bench serve \
  --base-url http://localhost:8000 \
  --dataset-name=custom \
  --model 1_run_model_in_vllm \
  --skip-tokenizer-init \
  --endpoint /pooling \
  --backend vllm-pooling \
  --percentile-metrics e2el \
  --metric-percentiles 25,75,99 \
  --num-prompts 100 \
  --request-rate 10 \
  --dataset-path ./dataset_path_input_india.jsonl
```

#Endpoint

#Name of the model

#skip tokenizer (EO models do not use it)

#server endpoint

#request handler

#additional end-to-end latency metric

#percentiles

#number of requests

#request rate (poisson default)

#datasetpath

vllm bench

```
INFO 05-15 15:45:05 [datasets.py:305] oversampled requests to reach 100
Starting initial single prompt test run...
Skipping endpoint ready check.
Starting main benchmark run...
Traffic request rate: 10.0
Burstiness factor: 1.0 (Poisson process)
Maximum request concurrency: None
100%|████████████████████████████████████████| 100/100 [00:10<00:00, 9.85it/s]
===== Serving Benchmark Result =====
Successful requests:          100
Failed requests:              0
Request rate configured (RPS): 10.00
Benchmark duration (s):       10.15
Total input tokens:           0
Request throughput (req/s):   9.85
-----End-to-end Latency-----
Mean E2EL (ms):               195.60
Median E2EL (ms):              167.47
P25 E2EL (ms):                 117.79
P75 E2EL (ms):                 262.21
P99 E2EL (ms):                 474.87
=====
```

GuideLLM sweep

```
guidellm benchmark \
  --target http://localhost:8000 \           #server endpoint
  --backend openai_http \                   #inference scenario (alternative: vllm_python)
  --model 1_run_model_in_vllm \            #model name
  --data dataset_path_input_india_guidellm.jsonl \ #dataset
  --request-format /pooling \              #format request for /pooling endpoint
  --data-column-mapper pooling_column_mapper \. #surface the correct key from dataset
  --max-requests 100 \                     # number of requests
  --output-path results.json               # save output to file
```


What to expect

Benchmarks									
[14:37:32]	100% synchronous	(complete)	Req:	11.2 req/s,	0.09s Lat,	1.0 Conc,	100 Comp,	0 Inc,	0 Err
			Tok:	0.0 gen/s,	0.0 tot/s,	0.0ms TTFT,	0.0ms ITL,	0 Prompt,	0 Gen
[14:37:43]	100% throughput@512	(complete)	Req:	18.8 req/s,	3.10s Lat,	58.4 Conc,	100 Comp,	0 Inc,	0 Err
			Tok:	0.0 gen/s,	0.0 tot/s,	0.0ms TTFT,	0.0ms ITL,	0 Prompt,	0 Gen
[14:37:50]	100% constant@12.17	(complete)	Req:	12.1 req/s,	0.13s Lat,	1.5 Conc,	100 Comp,	0 Inc,	0 Err
			Tok:	0.0 gen/s,	0.0 tot/s,	0.0ms TTFT,	0.0ms ITL,	0 Prompt,	0 Gen
[14:38:00]	100% constant@13.12	(complete)	Req:	13.1 req/s,	0.13s Lat,	1.7 Conc,	100 Comp,	0 Inc,	0 Err
			Tok:	0.0 gen/s,	0.0 tot/s,	0.0ms TTFT,	0.0ms ITL,	0 Prompt,	0 Gen
[14:38:09]	100% constant@14.07	(complete)	Req:	14.0 req/s,	0.12s Lat,	1.7 Conc,	100 Comp,	0 Inc,	0 Err
			Tok:	0.0 gen/s,	0.0 tot/s,	0.0ms TTFT,	0.0ms ITL,	0 Prompt,	0 Gen
[14:38:18]	100% constant@15.02	(complete)	Req:	14.9 req/s,	0.12s Lat,	1.8 Conc,	100 Comp,	0 Inc,	0 Err
			Tok:	0.0 gen/s,	0.0 tot/s,	0.0ms TTFT,	0.0ms ITL,	0 Prompt,	0 Gen
[14:38:26]	100% constant@15.96	(complete)	Req:	15.9 req/s,	0.12s Lat,	1.9 Conc,	100 Comp,	0 Inc,	0 Err
			Tok:	0.0 gen/s,	0.0 tot/s,	0.0ms TTFT,	0.0ms ITL,	0 Prompt,	0 Gen
[14:38:35]	100% constant@16.91	(complete)	Req:	16.6 req/s,	0.18s Lat,	2.9 Conc,	100 Comp,	0 Inc,	0 Err
			Tok:	0.0 gen/s,	0.0 tot/s,	0.0ms TTFT,	0.0ms ITL,	0 Prompt,	0 Gen
[14:38:42]	100% constant@17.86	(complete)	Req:	17.3 req/s,	0.21s Lat,	3.7 Conc,	100 Comp,	0 Inc,	0 Err
			Tok:	0.0 gen/s,	0.0 tot/s,	0.0ms TTFT,	0.0ms ITL,	0 Prompt,	0 Gen
[14:38:50]	100% constant@18.81	(complete)	Req:	17.8 req/s,	0.33s Lat,	6.0 Conc,	100 Comp,	0 Inc,	0 Err
			Tok:	0.0 gen/s,	0.0 tot/s,	0.0ms TTFT,	0.0ms ITL,	0 Prompt,	0 Gen

Generating...

GuideLLM sweep

Benchmarks			throughput	Mean latency	Mean concurrency	Requests completed		
			↓	↓	↓	↓		
[14:37:32]	100% synchronous	(complete)	Req: 11.2 req/s,	0.09s Lat,	1.0 Conc,	100 Comp,	0 Inc,	0 Err
			Tok: 0.0 gen/s,	0.0 tot/s,	0.0ms TTFT,	0.0ms ITL,	0 Prompt,	0 Gen
[14:37:43]	100% throughput@512	(complete)	Req: 18.8 req/s,	3.10s Lat,	58.4 Conc,	100 Comp,	0 Inc,	0 Err
			Tok: 0.0 gen/s,	0.0 tot/s,	0.0ms TTFT,	0.0ms ITL,	0 Prompt,	0 Gen
[14:37:50]	100% constant@12.17	(complete)	Req: 12.1 req/s,	0.13s Lat,	1.5 Conc,	100 Comp,	0 Inc,	0 Err
			Tok: 0.0 gen/s,	0.0 tot/s,	0.0ms TTFT,	0.0ms ITL,	0 Prompt,	0 Gen
[14:38:00]	100% constant@13.12	(complete)	Req: 13.1 req/s,	0.13s Lat,	1.7 Conc,	100 Comp,	0 Inc,	0 Err
			Tok: 0.0 gen/s,	0.0 tot/s,	0.0ms TTFT,	0.0ms ITL,	0 Prompt,	0 Gen
[14:38:09]	100% constant@14.07	(complete)	Req: 14.0 req/s,	0.12s Lat,	1.7 Conc,	100 Comp,	0 Inc,	0 Err
			Tok: 0.0 gen/s,	0.0 tot/s,	0.0ms TTFT,	0.0ms ITL,	0 Prompt,	0 Gen
[14:38:18]	100% constant@15.02	(complete)	Req: 14.9 req/s,	0.12s Lat,	1.8 Conc,	100 Comp,	0 Inc,	0 Err
			Tok: 0.0 gen/s,	0.0 tot/s,	0.0ms TTFT,	0.0ms ITL,	0 Prompt,	0 Gen
[14:38:26]	100% constant@15.96	(complete)	Req: 15.9 req/s,	0.12s Lat,	1.9 Conc,	100 Comp,	0 Inc,	0 Err
			Tok: 0.0 gen/s,	0.0 tot/s,	0.0ms TTFT,	0.0ms ITL,	0 Prompt,	0 Gen
[14:38:35]	100% constant@16.91	(complete)	Req: 16.6 req/s,	0.18s Lat,	2.9 Conc,	100 Comp,	0 Inc,	0 Err
			Tok: 0.0 gen/s,	0.0 tot/s,	0.0ms TTFT,	0.0ms ITL,	0 Prompt,	0 Gen
[14:38:42]	100% constant@17.86	(complete)	Req: 17.3 req/s,	0.21s Lat,	3.7 Conc,	100 Comp,	0 Inc,	0 Err
			Tok: 0.0 gen/s,	0.0 tot/s,	0.0ms TTFT,	0.0ms ITL,	0 Prompt,	0 Gen
[14:38:50]	100% constant@18.81	(complete)	Req: 17.8 req/s,	0.33s Lat,	6.0 Conc,	100 Comp,	0 Inc,	0 Err
			Tok: 0.0 gen/s,	0.0 tot/s,	0.0ms TTFT,	0.0ms ITL,	0 Prompt,	0 Gen

Generating...

GuideLLM sweep

Rates			throughput	Mean latency	Mean concurrency	Requests completed		
			↓	↓	↓	↓		
Benchmarks								
[14:37:32]	100% synchronous	(complete)	Req: 11.2 req/s,	0.09s Lat,	1.0 Conc,	100 Comp,	0 Inc,	0 Err
			Tok: 0.0 gen/s,	0.0 tot/s,	0.0ms TTFT,	0.0ms ITL,	0 Prompt,	0 Gen
[14:37:43]	100% throughput@512	(complete)	Req: 18.8 req/s,	3.10s Lat,	58.4 Conc,	100 Comp,	0 Inc,	0 Err
			Tok: 0.0 gen/s,	0.0 tot/s,	0.0ms TTFT,	0.0ms ITL,	0 Prompt,	0 Gen
[14:37:50]	100% constant@12.17	(complete)	Req: 12.1 req/s,	0.13s Lat,	1.5 Conc,	100 Comp,	0 Inc,	0 Err
			Tok: 0.0 gen/s,	0.0 tot/s,	0.0ms TTFT,	0.0ms ITL,	0 Prompt,	0 Gen
[14:38:00]	100% constant@13.12	(complete)	Req: 13.1 req/s,	0.13s Lat,	1.7 Conc,	100 Comp,	0 Inc,	0 Err
			Tok: 0.0 gen/s,	0.0 tot/s,	0.0ms TTFT,	0.0ms ITL,	0 Prompt,	0 Gen
[14:38:09]	100% constant@14.07	(complete)	Req: 14.0 req/s,	0.12s Lat,	1.7 Conc,	100 Comp,	0 Inc,	0 Err
			Tok: 0.0 gen/s,	0.0 tot/s,	0.0ms TTFT,	0.0ms ITL,	0 Prompt,	0 Gen
[14:38:18]	100% constant@15.02	(complete)	Req: 14.9 req/s,	0.12s Lat,	1.8 Conc,	100 Comp,	0 Inc,	0 Err
			Tok: 0.0 gen/s,	0.0 tot/s,	0.0ms TTFT,	0.0ms ITL,	0 Prompt,	0 Gen
[14:38:26]	100% constant@15.96	(complete)	Req: 15.9 req/s,	0.12s Lat,	1.9 Conc,	100 Comp,	0 Inc,	0 Err
			Tok: 0.0 gen/s,	0.0 tot/s,	0.0ms TTFT,	0.0ms ITL,	0 Prompt,	0 Gen
[14:38:35]	100% constant@16.91	(complete)	Req: 16.6 req/s,	0.18s Lat,	2.9 Conc,	100 Comp,	0 Inc,	0 Err
			Tok: 0.0 gen/s,	0.0 tot/s,	0.0ms TTFT,	0.0ms ITL,	0 Prompt,	0 Gen
[14:38:42]	100% constant@17.86	(complete)	Req: 17.3 req/s,	0.21s Lat,	3.7 Conc,	100 Comp,	0 Inc,	0 Err
			Tok: 0.0 gen/s,	0.0 tot/s,	0.0ms TTFT,	0.0ms ITL,	0 Prompt,	0 Gen
[14:38:50]	100% constant@18.81	(complete)	Req: 17.8 req/s,	0.33s Lat,	6.0 Conc,	100 Comp,	0 Inc,	0 Err
			Tok: 0.0 gen/s,	0.0 tot/s,	0.0ms TTFT,	0.0ms ITL,	0 Prompt,	0 Gen

Generating...

Benchmarking Time!

(follow [2_vllm_benchmarking/README.md](#))

Customize your benchmark with GuideLLM profiles

Profile	Goal	Best for
sweep	rate exploration	finding limits
throughput	max load	peak throughput
concurrent	# users	user-level latency
constant	fixed RPS	steady-state
poisson	randomized RPS	realistic traffic
synchronous	sequential	baseline latency

Recap so far

We learned how to:

- Prepare a TerraTorch model and serve it with vLLM
- Benchmark models with vLLM bench and GuideLLM
- Find the sweet spot between Throughput and Latency with GuideLLM


From Earth to Cloud

Scaling vLLM in Kubernetes

Works on my machine...
I swear!

Let's recap:

1. We logged onto a managed JupyterLab instance (on a virtual machine hosted on AWS)
- ~~2. We installed vLLM, TerraTorch, TIMM~~
3. Started vLLM and used it to serve Prithvi locally




Works on my machine...
I swear!

Let's recap:

1. We logged onto a managed JupyterLab instance (on a virtual machine hosted on AWS)
2. We installed vLLM, TerraTorch, TIMM
3. Started vLLM and used it to serve Prithvi locally

This is:

- *manual*
- *time consuming*
- *not reproducible*



Works on my machine...
And yours too!


Let's recap:

1. We logged onto a managed JupyterLab instance (on a virtual machine hosted on AWS)
2. We installed vLLM, TerraTorch, TIMM
3. Started vLLM and used it to serve Prithvi locally

This is:

- *manual*
- *time consuming*
- *not reproducible*

Instead of repeating setup steps every time, we can build one portable artifact: a ***container image***.



Works on my machine...
And yours too!

Let's recap:

1. We logged onto a managed JupyterLab instance (on a virtual machine hosted on AWS)
- ~~2. We installed vLLM, TerraTorch, TIMM~~
3. Started vLLM and used it to serve Prithvi locally

This is:

- *manual*
- *time consuming*
- *not reproducible*

Instead of repeating setup steps every time, we can build one portable artifact: a ***container image***.


```
# Use the official vLLM image with CUDA support as the base
# (already packages Python 3.12, uv, vLLM, CUDA toolkit, TIMM, and more)
FROM vllm/vllm-openai:v0.18.0
```

```
# Install terratorch with vLLM support at version 1.2.5
RUN pip3 install terratorch[vllm]==1.2.5
```

Containers need a ship!

Container runtimes like Docker and Podman can run container images as isolated environments where your app runs the same way every time.

With containers we have solved reproducibility issues and have removed manual steps, but ***on their own they aren't production ready.***



Containers need a ship!


Container runtimes like Docker and Podman can run container images as isolated environments where your app runs the same way every time.

With containers we have solved reproducibility issues and have removed manual steps, but *on their own they aren't production ready*.

We need a tool that can:

- Schedule containers across multiple nodes
- Restart them if they crash
- Add/Remove containers as demand changes
- Expose endpoints to the network


*The most widely used tool for this is **Kubernetes**.*



Kubernetes in 60 seconds (only what we need)

Kubernetes introduces a lot of terminology, let's focus on only a subset:


- **Node**: a single machine
- **Pod**: the smallest execution unit in Kubernetes. It's typically a single container.
- **Deployment**: object that manages Pods, keeping N of them running, handling crashes and updates
- **Service**: a stable, internal network address for a set of pods
- **Ingress/Gateway/Route**: a way to make Services reachable outside of the cluster



Kubernetes in 60 seconds (only what we need)

Kubernetes introduces a lot of terminology, let's focus on only a subset:

- **Node:** a single machine
- **Pod:** the smallest execution unit in Kubernetes. It's typically a single container.
- **Deployment:** object that manages Pods, keeping N of them running, handling crashes and updates
- **Service:** a stable, internal network address for a set of pods
- **Ingress/Gateway/Route:** a way to make Services reachable outside of the cluster




From containers...

So, we can:

- Package our execution environment in a container image.
- Run it via a container runtime.
- Use Kubernetes to get production-grade orchestration.

I didn't mention anything specific to vLLM or model serving...



From containers...

So, we can:

- Package our execution environment in a container image.
- Run it via a container runtime.
- Use Kubernetes to get production-grade orchestration.

I didn't mention anything specific to vLLM or model serving...

Kubernetes is designed to be generic and extensible. Projects like **KServe** extend it adding capabilities like model serving on top of Kubernetes.



From containers...
... to Inference Services

KServe introduces the
InferenceService abstraction.

In it we declare:

- What we want to serve
- What resources it should get
- How it should scale
- How it should be exposed

*KServe will create the
underlying Deployments,
Services, and Routes for us.*




From containers...
... to Inference Services

KServe introduces the
InferenceService abstraction.

In it we declare:

- What we want to serve
- What resources it should get
- How it should scale
- How it should be exposed

*KServe will create the
underlying Deployments,
Services, and Routes for us.*



From containers...
... to Inference Services


KServe introduces the
InferenceService abstraction.

In it we declare:

- What we want to serve
- What resources it should get
- How it should scale
- How it should be exposed

*KServe will create the
underlying Deployments,
Services, and Routes for us.*

So... how does it look like in
practice?



From containers...
... to Inference Services

KServe introduces the *InferenceService* abstraction.

In it we declare:

- What we want to serve
- What resources it should get
- How it should scale
- How it should be exposed

KServe will create the underlying Deployments, Services, and Routes for us.

So... how does it look like in practice?

```
apiVersion: serving.kserve.io/v1beta1
kind: InferenceService
metadata:
  name: prithvi-sen1floods11
annotations:
  serving.kserve.io/deploymentMode: "RawDeployment"
  serving.kserve.io/autoscalerClass: "keda"
  prometheus.io/scrape: "true"
  prometheus.io/path: "/metrics"
  prometheus.io/port: "8000"
  prometheus.io/scheme: "http"
labels:
  networking.kserve.io/visibility: exposed
spec:
  predictor:

    containers:
      - name: vllm
        image: quay.io/alessandropomponio_ibm/vllm:v0.18.0-tt.v1.2.5
        command: ["vllm", "serve"]
        args:
          - "ibm-nasa-geospatial/Prithvi-EO-2.0-300M-TL-Sen1Floods11"
          - "--skip-tokenizer-init"
          - "--enable-mm-embeds"
          - "--io-processor-plugin"
          - "terratorch_segmentation"
          - "--max-num-seqs"
          - "1"
          - "--enforce-eager"

    resources:
      limits:
        cpu: "64"
        memory: 64G
        nvidia.com/gpu: "1"
      requests:
        cpu: "32"
        memory: 64G
        nvidia.com/gpu: "1"

  minReplicas: 1
  maxReplicas: 5
  autoScaling:
    metrics:
      - type: External
        external:
          metric:
            backend: "prometheus"
            serverAddress: "https://thanos-querier.openshift-monitoring.svc:9091"
            query: "sum(rate(vllm:request_success_total[1m]))"
          target:
            type: Value
            value: 5
```

From containers... ... to Inference Services

KServe introduces the *InferenceService* abstraction.

In it we declare:

- What we want to serve
- What resources it should get
- How it should scale
- How it should be exposed

KServe will create the underlying Deployments, Services, and Routes for us.

So... how does it look like in practice?

```
apiVersion: serving.kserve.io/v1beta1
kind: InferenceService
metadata:
  name: prithvi-sen1floods11
annotations:
  serving.kserve.io/deploymentMode: "RawDeployment"
  serving.kserve.io/autoscalerClass: "keda"
  prometheus.io/scrape: "true"
  prometheus.io/path: "/metrics"
  prometheus.io/port: "8000"
  prometheus.io/scheme: "http"
labels:
  networking.kserve.io/visibility: exposed
spec:
  predictor:

containers:
  - name: vllm
    image: quay.io/alessandropomponio_ibm/vllm:v0.18.0-tt.v1.2.5
    command: ["vllm", "serve"]
    args:
      - "ibm-nasa-geospatial/Prithvi-EO-2.0-300M-TL-Sen1Floods11"
      - "--skip-tokenizer-init"
      - "--enable-mm-embeds"
      - "--io-processor-plugin"
      - "terratorch_segmentation"
      - "--max-num-seqs"
      - "1"
      - "--enforce-eager"

resources:
  limits:
    cpu: "64"
    memory: 64G
    nvidia.com/gpu: "1"
  requests:
    cpu: "32"
    memory: 64G
    nvidia.com/gpu: "1"

minReplicas: 1
maxReplicas: 5
autoScaling:
  metrics:
    - type: External
      external:
        metric:
          backend: "prometheus"
          serverAddress: "https://thanos-querier.openshift-monitoring.svc:9091"
          query: "sum(rate(vllm:request_success_total[1m]))"
        target:
          type: Value
          value: 5
```


From containers... ... to Inference Services

KServe introduces the *InferenceService* abstraction.

In it we declare:

- What we want to serve
- What resources it should get
- How it should scale
- How it should be exposed

KServe will create the underlying Deployments, Services, and Routes for us.

So... how does it look like in practice?

```
apiVersion: serving.kserve.io/v1beta1
kind: InferenceService
metadata:
  name: prithvi-sen1floods11
annotations:
  serving.kserve.io/deploymentMode: "RawDeployment"
  serving.kserve.io/autoscalerClass: "keda"
  prometheus.io/scrape: "true"
  prometheus.io/path: "/metrics"
  prometheus.io/port: "8000"
  prometheus.io/scheme: "http"
labels:
  networking.kserve.io/visibility: exposed
spec:
  predictor:

    containers:
      - name: vllm
        image: quay.io/alessandropomponio_ibm/vllm:v0.18.0-tt.v1.2.5
        command: ["vllm", "serve"]
        args:
          - "ibm-nasa-geospatial/Prithvi-EO-2.0-300M-TL-Sen1Floods11"
          - "--skip-tokenizer-init"
          - "--enable-mm-embeds"
          - "--io-processor-plugin"
          - "terratorch_segmentation"
          - "--max-num-seqs"
          - "1"
          - "--enforce-eager"

        resources:
          limits:
            cpu: "64"
            memory: 64G
            nvidia.com/gpu: "1"
          requests:
            cpu: "32"
            memory: 64G
            nvidia.com/gpu: "1"

    minReplicas: 1
    maxReplicas: 5
    autoScaling:
      metrics:
        - type: External
          external:
            metric:
              backend: "prometheus"
              serverAddress: "https://thanos-querier.openshift-monitoring.svc:9091"
              query: "sum(rate(vllm:request_success_total[1m]))"
            target:
              type: Value
              value: 5
```


From containers...
... to Inference Services

KServe introduces the *InferenceService* abstraction.

In it we declare:

- What we want to serve
- What resources it should get
- How it should scale
- How it should be exposed

KServe will create the underlying Deployments, Services, and Routes for us.

So... how does it look like in practice?

```
apiVersion: serving.kserve.io/v1beta1
kind: InferenceService
metadata:
  name: prithvi-sen1floods11
annotations:
  serving.kserve.io/deploymentMode: "RawDeployment"
  serving.kserve.io/autoscalerClass: "keda"
  prometheus.io/scrape: "true"
  prometheus.io/path: "/metrics"
  prometheus.io/port: "8000"
  prometheus.io/scheme: "http"
labels:
  networking.kserve.io/visibility: exposed
spec:
  predictor:

    containers:
      - name: vllm
        image: quay.io/alessandropomponio_ibm/vllm:v0.18.0-tt.v1.2.5
        command: ["vllm", "serve"]
        args:
          - "ibm-nasa-geospatial/Prithvi-EO-2.0-300M-TL-Sen1Floods11"
          - "--skip-tokenizer-init"
          - "--enable-mm-embeds"
          - "--io-processor-plugin"
          - "terratorch_segmentation"
          - "--max-num-seqs"
          - "1"
          - "--enforce-eager"

    resources:
      limits:
        cpu: "64"
        memory: 64G
        nvidia.com/gpu: "1"
      requests:
        cpu: "32"
        memory: 64G
        nvidia.com/gpu: "1"

  minReplicas: 1
  maxReplicas: 5
  autoScaling:
    metrics:
      - type: External
        external:
          metric:
            backend: "prometheus"
            serverAddress: "https://thanos-querier.openshift-monitoring.svc:9091"
            query: "sum(rate(vllm:request_success_total[1m]))"
          target:
            type: Value
            value: 5
```

From containers...
... to Inference Services

KServe introduces the *InferenceService* abstraction.

In it we declare:

- What we want to serve
- What resources it should get
- How it should scale
- How it should be exposed

KServe will create the underlying Deployments, Services, and Routes for us.

So... how does it look like in practice?

```
apiVersion: serving.kserve.io/v1beta1
kind: InferenceService
metadata:
  name: prithvi-sen1floods11
  annotations:
    serving.kserve.io/deploymentMode: "RawDeployment"
    serving.kserve.io/autoscalerClass: "keda"
    prometheus.io/scrape: "true"
    prometheus.io/path: "/metrics"
    prometheus.io/port: "8000"
    prometheus.io/scheme: "http"
  labels:
    networking.kserve.io/visibility: exposed
spec:
  predictor:

    containers:
      - name: vllm
        image: quay.io/alessandropomponio_ibm/vllm:v0.18.0-tt.v1.2.5
        command: ["vllm", "serve"]
        args:
          - "ibm-nasa-geospatial/Prithvi-EO-2.0-300M-TL-Sen1Floods11"
          - "--skip-tokenizer-init"
          - "--enable-mm-embeds"
          - "--io-processor-plugin"
          - "terratorch_segmentation"
          - "--max-num-seqs"
          - "1"
          - "--enforce-eager"

    resources:
      limits:
        cpu: "64"
        memory: 64G
        nvidia.com/gpu: "1"
      requests:
        cpu: "32"
        memory: 64G
        nvidia.com/gpu: "1"

  minReplicas: 1
  maxReplicas: 5
  autoScaling:
    metrics:
      - type: External
        external:
          metric:
            backend: "prometheus"
            serverAddress: "https://thanos-querier.openshift-monitoring.svc:9091"
            query: "sum(rate(vllm:request_success_total[1m]))"
          target:
            type: Value
            value: 5
```

From containers...
... to Inference Services

KServe introduces the *InferenceService* abstraction.

In it we declare:

- What we want to serve
- What resources it should get
- How it should scale
- How it should be exposed

KServe will create the underlying Deployments, Services, and Routes for us.

So... how does it look like in practice?

```
apiVersion: serving.kserve.io/v1beta1
kind: InferenceService
metadata:
  name: prithvi-sen1floods11
annotations:
  serving.kserve.io/deploymentMode: "RawDeployment"
  serving.kserve.io/autoscalerClass: "keda"
  prometheus.io/scrape: "true"
  prometheus.io/path: "/metrics"
  prometheus.io/port: "8000"
  prometheus.io/scheme: "http"
labels:
  networking.kserve.io/visibility: exposed
spec:
  predictor:

  containers:
    - name: vllm
      image: quay.io/alessandropomponio_ibm/vllm:v0.18.0-tt.v1.2.5
      command: ["vllm", "serve"]
      args:
        - "ibm-nasa-geospatial/Prithvi-EO-2.0-300M-TL-Sen1Floods11"
        - "--skip-tokenizer-init"
        - "--enable-mm-embeds"
        - "--io-processor-plugin"
        - "terratorch_segmentation"
        - "--max-num-seqs"
        - "1"
        - "--enforce-eager"

  resources:
    limits:
      cpu: "64"
      memory: 64G
      nvidia.com/gpu: "1"
    requests:
      cpu: "32"
      memory: 64G
      nvidia.com/gpu: "1"

  minReplicas: 1
  maxReplicas: 5
  autoScaling:
    metrics:
      - type: External
        external:
          metric:
            backend: "prometheus"
            serverAddress: "https://thanos-querier.openshift-monitoring.svc:9091"
            query: "sum(rate(vllm:request_success_total[1m]))"
          target:
            type: Value
            value: 5
```



Demo: KServe + KEDA scaling
in action

Let's see it in practice with a
simple *InferenceService* and
see event-driven autoscaling in
action by generating some load
with *vllm bench*



Thank you

*Check out the GitHub repo
using the QR Code.*



© Copyright IBM Corporation 2026

IBM and the IBM logo are trademarks of IBM Corporation, registered in many jurisdictions worldwide. Other product and service names might be trademarks of IBM or other companies. A current list of IBM trademarks is available on ibm.com/legal/copytrade.

This document is current as of the initial date of publication and may be changed by IBM at any time.

THE INFORMATION IN THIS DOCUMENT IS PROVIDED "AS IS" WITHOUT ANY WARRANTY, EXPRESS OR IMPLIED, INCLUDING WITHOUT ANY WARRANTIES OF MERCHANTABILITY, FITNESS FOR A PARTICULAR PURPOSE AND ANY WARRANTY OR CONDITION OF NON-INFRINGEMENT.

Examples presented as illustrative only. Actual results will vary based on client configurations and conditions and, therefore, generally expected results cannot be provided.

Not all offerings are available in every country in which IBM operates.

Statements regarding IBM's future direction and intent are subject to change or withdrawal without notice, and represent goals and objectives only.